


Governed AI Agents

Putting Hermes Agent and OpenClaw to work safely inside your business.



Agents that get work done inside your walls, with you in control.

AI agents go beyond answering questions. Open-source agents such as Hermes Agent and OpenClaw can read files, research, draft documents, run scheduled tasks and act across tools, taking instructions from the messaging apps people already use.

That capability is exactly why agents need governance. Logic Sonata deploys them inside a governed runtime: private models on your appliance, sandboxed tools with allow-listed access, credentials held in a vault, human approval for external or irreversible actions and a complete audit trail.

Key takeaways

- ✓ Real work, not just chat: reports, research, drafting and follow-ups handled end to end.
- ✓ Bounded by design: agents only touch the files, systems and network access each task needs.
- ✓ People stay in charge: sensitive steps wait for approval and everything is recorded.

This paper explains how the Private AI Agent Platform works, the controls that keep data private, where it delivers value first, and how to evaluate and roll it out.

Agents are powerful. Ungoverned agents are a risk.

Open-source agents can read files, browse the web, run commands and send messages. That is exactly what makes them useful and exactly why they must not run with unchecked access to company systems or public AI services. The value comes from pairing capable agents with strict governance.


- ✓ Teams spend hours each week compiling reports from exports and spreadsheets.
- ✓ Follow-ups with suppliers and customers fall through the cracks.
- ✓ Open-source agents are powerful but risky when given broad access or public AI models.
- ✓ Without logs, nobody can reconstruct what an automated process did or why.



Illustration · The governed agent loop: every step bounded, approved and recorded.

From request to result, inside your walls.

Hermes Agent from Nous Research or OpenClaw, both open-source, self-hosted agent frameworks with memory, skills, scheduling and messaging integrations. Logic Sonata selects the framework per use case and adds the governed runtime.



01 Trigger

A message, a schedule or an event starts a task.

02 Plan

The agent breaks the task into steps using a private model on the appliance.

03 Act in the sandbox

Approved tools run in isolation with only the access that task needs.

04 Ask before it matters

Sensitive steps pause for human approval in chat or the console.

05 Deliver and learn

Results are delivered, logged and captured as reusable skills.




Illustration • The agent plans, works in its sandbox and asks before sending anything out.

What the Private AI Agent Platform does.

Work from chat

Ask the agent in messaging apps such as Telegram, WhatsApp, Slack or Discord, or from a web console.

Memory and skills

Agents remember context across sessions and turn solved problems into reusable, readable skills.

Scheduled tasks

Daily reports, weekly summaries and recurring checks, described in plain language.

Sandboxed tools

Files, spreadsheets, browser and scripts run in an isolated environment with allow-listed access.

Human approval

Sending email, updating records or spending money waits for a named approver.

Full audit trail

Every plan, tool call, approval and result is recorded and reviewable.

Built on proven open foundations

Hermes Agent from Nous Research or OpenClaw, both open-source, self-hosted agent frameworks with memory, skills, scheduling and messaging integrations. Logic Sonata selects the framework per use case and adds the governed runtime.

Controls built in, not bolted on.

Every Logic Sonata system runs on infrastructure you control. Data boundaries, access and audit are designed in from the start rather than added later.

- ✓ Agents run on private models, never public AI services
- ✓ Isolated sandbox with allow-listed tools and network access
- ✓ Credentials held in a vault, never exposed to the model
- ✓ Approval gates for external and irreversible actions
- ✓ Complete, reviewable audit trail
- ✓ Instant pause and kill switch

Data boundary

Prompts, documents, images and results are processed on your appliance or dedicated private infrastructure. They are not sent to public AI services and are never used to train third-party models.

Work your agents can take on.

Skills are readable, reusable instructions the agent builds and improves. These are typical first skills for manufacturing and trading companies.

Every morning

Read yesterday's production export, calculate output, efficiency and defects by line, and post a summary to the operations channel.

On request

Collect the buyer's tech pack, past quotations and material prices into a first-draft RFQ response for review.

Every Friday

List purchase orders overdue by more than three days and draft supplier reminders for approval.

After meetings

Summarise the notes, extract decisions and owners, and post them back to the team.

Where it pays off first.

Daily production and sales reports

Pull exports from ERP or MES, calculate KPIs and deliver a summary before the morning meeting.

RFQ and tender preparation

Collect specifications, past quotes and requirements into a first draft for review.

Supplier follow-ups

Track open purchase orders and draft chasers for approval.

Inbox and document triage

Classify incoming requests, extract key data and route them to the right team.

Meeting notes and actions

Summaries, decisions and owners, posted back to the team channel.

Market and competitor watch

Research through approved sources, summarised weekly.

Measuring the pilot

We agree success measures before the pilot starts. Typical measures include:

- ✓ Hours saved per week on the first automated reports
- ✓ On-time rate of supplier follow-ups
- ✓ Share of agent actions completed without rework
- ✓ Approval turnaround time

Deploy where your data needs to live.

On-premise

Dedicated hardware inside your building for the most sensitive data. Nothing leaves your network.

Hosted private

Dedicated capacity on managed private GPU infrastructure, isolated from public AI services. No hardware purchase.

Hybrid

Sensitive workloads on-site, everything else hosted, under one set of policies and one audit trail.

Hardware, sized to the workload

Logic Sonata is hardware-neutral. Smaller teams often start with a compact AI workstation such as NVIDIA DGX Spark or an AMD Ryzen AI Max+ system, both with up to 128 GB of unified memory in a desktop-sized unit. Larger rollouts use multi-GPU servers or a private cloud cluster.

A proven rollout path

Assess

Review current AI use, data exposure and priority use cases. Deliver a 30/60/90-day roadmap.

Pilot

Deploy one working use case with real users and agreed success measures.

Roll out

Extend to more teams with governance, integration and role-based training.

Manage

Monitoring, model updates, knowledge refreshes and quarterly reviews.

One partner is accountable for the whole stack: hardware, software and models, your company knowledge, access controls, implementation, training and ongoing support.

Questions to ask any vendor.

Use this checklist when comparing options, including ours.

- ☐ Which model powers the agent, and where does it run?
- ☐ What tools and systems can the agent access, and how is that scoped per task?
- ☐ Where are credentials stored, and can the model ever see them?
- ☐ Which actions require human approval, and who approves them?
- ☐ Can we pause or stop an agent instantly?
- ☐ Is every plan, tool call and result recorded for review?

Frequently asked questions

Hermes Agent or OpenClaw: which do we get?

Both are capable open-source agents. We recommend one based on your channels, tasks and security requirements, and deploy it inside the same governed runtime.

Can an agent act without permission?

Only within the limits you set. Read-only and internal steps can run automatically; anything external or irreversible waits for approval.

Which systems can agents connect to?

Files, spreadsheets, email, messaging apps and business systems with an API or export. Each connection is allow-listed and scoped.

Private AI, built inside your walls.

Logic Sonata delivers secure private AI for businesses in Singapore, Vietnam, Indonesia, Malaysia and Thailand. Each solution combines hardware, software, company knowledge, access controls, implementation, training and ongoing support, deployed on-premise, hosted or hybrid.

Our product portfolio: Private Knowledge Assistant, Private Vision Intelligence, Private Coding Assistant, Private AI Agent Platform and Private Image Generation Studio.

Next step: a 30-minute consultation

We will show the Private AI Agent Platform working, discuss your data and requirements, and scope a pilot around your use case.

www.logicsonata.com/contact

sales@logicsonata.com

Product names mentioned are trademarks of their respective owners. Illustrations show representative interfaces and synthetic data. This document is for information only and does not constitute an offer.