


Vision AI on Your Own Cameras

Real-time operational insight without sending video to the cloud.



Your cameras, understood in plain language. On-site.

Factories, warehouses and stores already have cameras, but footage is usually reviewed only after something goes wrong. Traditional computer vision can automate some of that review, yet it needs a separately trained model for each question, which puts it out of reach for most mid-sized operations.

Vision-language models change the economics. They understand scenes described in everyday language, so supervisors can ask whether a station is working, whether bundles are queuing or whether a forklift lane is clear, and change the question at any time. Running these models on an appliance on-site keeps video and results inside the business.

Key takeaways

- ✓ One system, many questions: prompts replace custom model training for many use cases.
- ✓ Video never leaves the premises: analysis and storage happen on your appliance.
- ✓ Designed for responsible use: focus on processes and safety, with human review of findings.

This paper explains how the Private Vision Intelligence works, the controls that keep data private, where it delivers value first, and how to evaluate and roll it out.

Cameras record everything. Almost nobody watches.

Factories, warehouses and stores already have cameras, but footage is only reviewed after something goes wrong. Traditional computer vision needs a custom model for every question. A vision-language model understands scenes described in everyday language, so one system can answer many operational questions, live.


- ✓ Line imbalance and idle time are hard to see in real time, and manual time studies are infrequent.
- ✓ Safety and PPE compliance depend on supervisors being in the right place at the right time.
- ✓ Security teams cannot watch every camera, especially after hours.
- ✓ Cloud video analytics raise privacy concerns and ongoing bandwidth costs.



Illustration • Existing cameras over a production floor become a live source of operational insight.

From request to result, inside your walls.

NVIDIA Live VLM WebUI, an open-source real-time interface for vision-language models, with open models served on the appliance. Logic Sonata adds camera integration, prompt design, event logging, dashboards and governance.



01 Connect cameras

Add webcams or existing IP cameras. No new camera hardware is needed for most sites.

02 Describe what matters

Write prompts such as "Is the operator at station 4 working, waiting or away?"

03 Analyse live on-site

The appliance samples frames continuously and a vision-language model interprets each scene.

04 Act on the findings

Supervisors see live observations and shift summaries, and receive alerts for agreed events.

05 Refine and extend

Adjust prompts, add cameras and new questions as your team learns what is useful.



Illustration • Live observations, station status and shift timeline, generated on-site.

What the Private Vision Intelligence does.

Plain-language prompts

Describe what to look for in a sentence.
Change the question at any time without retraining a model.

Live camera streams

Browser webcams over WebRTC and IP cameras over RTSP feed the model continuously.

Runs on your appliance

Open vision-language models are served locally, so video and results stay on-site.

Real-time responses

Observations stream back within seconds, with latency and GPU utilisation shown alongside.

Structured events

Findings can be logged, summarised by shift and pushed to dashboards or alerts.

Human review

Observations are indicators for supervisors to review, not automated decisions about people.

Built on proven open foundations

NVIDIA Live VLM WebUI, an open-source real-time interface for vision-language models, with open models served on the appliance. Logic Sonata adds camera integration, prompt design, event logging, dashboards and governance.

Controls built in, not bolted on.

Every Logic Sonata system runs on infrastructure you control. Data boundaries, access and audit are designed in from the start rather than added later.

- ✓ Video and results processed and stored on-site only
- ✓ Configurable retention, with no footage kept by default
- ✓ Focus on processes and safety, not individual scoring
- ✓ Access to live views and history limited by role
- ✓ Deployment guidance aligned with local data protection laws
- ✓ Signage and worker consultation recommended before go-live

Data boundary

Prompts, documents, images and results are processed on your appliance or dedicated private infrastructure. They are not sent to public AI services and are never used to train third-party models.

Ask the line what is happening.

Instead of training a model per question, supervisors describe what they care about. The model answers continuously, and results roll up into shift summaries.

Prompt

Is the operator at station 4 sewing, waiting for work or away from the station?

Prompt

Are bundles piling up between stations 3 and 4? Describe the queue.

Prompt

Is anyone in the marked forklift lane without a high-visibility vest?

Prompt

After 22:00, report any person or vehicle near the loading bay door.

Where it pays off first.

Time and motion on the line

Detect idle stations, waiting for materials and unnecessary movement to balance lines and remove bottlenecks.

Quality spot checks

A second pair of eyes for visible defects such as stains, broken stitches or shade variation.

Safety and PPE

Flag missing helmets or vests, blocked exits and people too close to moving equipment.

Security monitoring

Activity after hours, loitering near restricted doors and unattended items, described in plain language.

Warehouse and logistics

Dock occupancy, loading progress and unsafe pallet stacking, shift by shift.

Retail operations

Queue length, empty shelves and fitting-room waits, so staff go where they are needed.

Measuring the pilot

We agree success measures before the pilot starts. Typical measures include:

- ✓ Idle and waiting time per station, before and after
- ✓ Bottleneck duration on the target line
- ✓ Safety observations per shift and time to resolution
- ✓ Precision of observations against reviewed footage

Deploy where your data needs to live.

On-premise

Dedicated hardware inside your building for the most sensitive data. Nothing leaves your network.

Hosted private

Dedicated capacity on managed private GPU infrastructure, isolated from public AI services. No hardware purchase.

Hybrid

Sensitive workloads on-site, everything else hosted, under one set of policies and one audit trail.

Hardware, sized to the workload

Logic Sonata is hardware-neutral. Smaller teams often start with a compact AI workstation such as NVIDIA DGX Spark or an AMD Ryzen AI Max+ system, both with up to 128 GB of unified memory in a desktop-sized unit. Larger rollouts use multi-GPU servers or a private cloud cluster.

A proven rollout path

Assess

Review current AI use, data exposure and priority use cases. Deliver a 30/60/90-day roadmap.

Pilot

Deploy one working use case with real users and agreed success measures.

Roll out

Extend to more teams with governance, integration and role-based training.

Manage

Monitoring, model updates, knowledge refreshes and quarterly reviews.

One partner is accountable for the whole stack: hardware, software and models, your company knowledge, access controls, implementation, training and ongoing support.

Questions to ask any vendor.

Use this checklist when comparing options, including ours.

- ☐ Can we connect our existing IP cameras, or do we need new hardware?
- ☐ Where are frames processed and stored, and for how long?
- ☐ Can supervisors change what the system looks for without retraining a model?
- ☐ How are findings reviewed before any action is taken?
- ☐ How does the deployment align with local data protection law and our works council or unions?
- ☐ What accuracy did the pilot achieve on our own footage?
- ☐ Who can see live views and history, and is access logged?

Frequently asked questions

Do we need new cameras?

Usually not. Webcams and most IP cameras that provide an RTSP stream can be connected. We check your cameras during the assessment.

Is this employee surveillance?

The system is designed to improve processes, safety and security. We recommend focusing prompts on stations and events rather than individuals, informing staff, and following local data protection law. Findings are for human review.

How accurate is it?

Vision-language models are strong at describing scenes but can make mistakes, so observations are treated as indicators. We tune prompts and camera placement during the pilot and measure accuracy against real footage.

Private AI, built inside your walls.

Logic Sonata delivers secure private AI for businesses in Singapore, Vietnam, Indonesia, Malaysia and Thailand. Each solution combines hardware, software, company knowledge, access controls, implementation, training and ongoing support, deployed on-premise, hosted or hybrid.

Our product portfolio: Private Knowledge Assistant, Private Vision Intelligence, Private Coding Assistant, Private AI Agent Platform and Private Image Generation Studio.

Next step: a 30-minute consultation

We will show the Private Vision Intelligence working, discuss your data and requirements, and scope a pilot around your use case.

www.logicsonata.com/contact

sales@logicsonata.com

Product names mentioned are trademarks of their respective owners. Illustrations show representative interfaces and synthetic data. This document is for information only and does not constitute an offer.